

Getting Started with LatinCy

Patrick J. Burns

2024-08-20

Table of contents

Preface	4
Key links	4
Acknowledgements	4
Abbreviations	5
I Front Matter	6
II Model/Pipeline Basics	7
1 Installing LatinCy models	8
1.1 Installing spaCy	8
1.2 Installing the LatinCy models	8
1.3 Installing additional packages	8
2 Loading LatinCy models	10
2.1 Loading LatinCy models with spaCy	10
2.2 Creating a spaCy doc from a string	10
III NLP Tasks	12
3 Sentence segmentation	13
3.1 Sentence segmentation with LatinCy	13
References	16
4 Word Tokenization	17
4.1 Word tokenization with LatinCy	17
4.1.1 Token attributes and methods related to tokenization	19
4.1.2 Customization of the spaCy tokenizer in LatinCy	22
References	25

5	Lemmatization	26
5.1	Lemmatization with LatinCy	26
	References	30
6	POS Tagging	31
6.1	POS Tagging with LatinCy	31
	References	36
IV	spaCy Tasks	37
	Appendices	38
	Open LatinCy Projects	38
	Model-based enclitic splitting	38
	References	39

Preface

“Getting Started with LatinCy” is an always-a-work-in-progress combination of documentation and demo notebooks for working with the LatinCy models on a variety of Latin text analysis and NLP tasks.

latinCy

Key links

Models: <https://huggingface.co/latincy>

Universe: <https://spacy.io/universe/project/latincy>

Preprint: <https://arxiv.org/abs/2305.04365>

This book has been written using Jupyter notebooks which have then been collated with Quarto. To learn more about Quarto books visit <https://quarto.org/docs/books>.

Acknowledgements

Thank you to all of the readers who have submitted corrections via GitHub PRs: [@sjhuskey](#).

Abbreviations

Where possible, I will include references to standard NLP works using the following abbreviations:

NLTK Bird, S., Klein, E., and Loper, E. 2015. *Natural Language Processing with Python: Analyzing Text with the Natural Language Toolkit*. 2nd Edition. <https://www.nltk.org/book/>. (Bird, Klein, and Loper 2015)

SLP Jurafsky, D., and Martin, J.H. 2020. *Speech and Language Processing*. 3rd Edition, Draft. <https://web.stanford.edu/~jurafsky/slp3/>. (Jurafsky and Martin 2020)

Part I

Front Matter

Part II

Model/Pipeline Basics

1 Installing LatinCy models

1.1 Installing spaCy

The LatinCy models are designed to work with the spaCy natural language platform, so you will need to have this package installed before anything else. The following cell has the pip install command for spaCy. At the time of writing, the latest version available for spaCy compatible with LatinCy is v3.7.5.

NB: To run the cells below, uncomment the commands by removing the `#` at the beginning of the line. The exclamation point at the beginning of the line is shorthand for the `%system magic command` in Jupyter which can be used to run shell commands from within a notebook.

```
# !pip install -U spacy
```

1.2 Installing the LatinCy models

LatinCy models are currently available in three sizes: ‘sm’, ‘md’, and ‘lg’. We will use the different models throughout the tutorials, so let’s install all three now so that they are available for future chapters.

```
# ! pip install "la-core-web-sm @ https://huggingface.co/latincy/la_core_web_sm/resolve/main,  
# ! pip install "la-core-web-md @ https://huggingface.co/latincy/la_core_web_md/resolve/main,  
# ! pip install "la-core-web-lg @ https://huggingface.co/latincy/la_core_web_lg/resolve/main,  
# ! pip install "la-core-web-trf @ https://huggingface.co/latincy/la_core_web_trf/resolve/main
```

1.3 Installing additional packages

We will also use other packages throughout these tutorials. They are included here for your convenience, as well as in the `requirements.txt` file in the code repository for this Quarto book.


```
# !pip install pandas
# !pip install matplotlib
# !pip install seaborn
# !pip install scikit-learn
# !pip install tqdm
```

2 Loading LatinCy models

2.1 Loading LatinCy models with spaCy

```
# Imports

import spacy
from pprint import pprint

nlp = spacy.load('la_core_web_lg')
```

2.2 Creating a spaCy doc from a string

```
text = "Haec narrantur a poetis de Perseo."
doc = nlp(text)
print(doc)
```

Haec narrantur a poetis de Perseo.

```
print(type(doc))
```

<class 'spacy.tokens.doc.Doc'>

```
print(doc.__repr__())
```

Haec narrantur a poetis de Perseo.

```
print(doc.text)
```

Haec narrantur a poetis de Perseo.

```
print(type(doc.text))
```

```
<class 'str'>
```

Part III

NLP Tasks

3 Sentence segmentation

3.1 Sentence segmentation with LatinCy

Sentence segmentation is the task of splitting a text into sentences. For the LatinCy models, this is a task been trained using spaCy’s **senter** factory to terminate sentences at both strong and weak stops, following the example of Clayman (1981) (see also, Wake (1957), Janson (1964)), who writes: “If all stops are made equivalent, i.e. if no distinction is made between the strong stop, weak stop and interrogation mark, editorial differences will be kept to a minimum.”

Given a spaCy Doc, the **sents** attribute will produce a **generator** object with the sentence from that document as determined by the dependency parser. Each sentence is a **Span** object with the start and end token indices from the original Doc.

```
# Imports & setup

import spacy
from pprint import pprint
nlp = spacy.load('la_core_web_lg')
text = "Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius."
doc = nlp(text)
print(doc)
```

Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius.

```
sents = doc.sents
print(type(sents))
```

```
<class 'generator'>
```

Like all **Span** objects, the text from each sentence can be retrieved with the **text** attribute. For convenience below, we convert the generator to list so that we can iterate over it multiple times. Here are the three (3) sentences identified in the example text as well as an indication of the sentences’ type, i.e. `<class 'spacy.tokens.span.Span'>`.

```
sents = list(sents)

for i, sent in enumerate(sents, 1):
    print(f'{i}: {sent.text}')
```

```
1: Haec narrantur a poetis de Perseo.
2: Perseus filius erat Iovis, maximi deorum.
3: Avus eius Acrisius appellabatur.
```

```
sent = sents[0]
print(type(sent))
```

```
<class 'spacy.tokens.span.Span'>
```

Sentences have the same attributes/methods available to them as any span (listed in the next cell). Following are some attributes/methods that may be particularly relevant to working with sentences.

```
sent_methods = [item for item in dir(sent) if '_' not in item]
pprint(sent_methods)
```

```
['conjuncts',
 'doc',
 'end',
 'ents',
 'id',
 'label',
 'lefts',
 'rights',
 'root',
 'sent',
 'sentiment',
 'sents',
 'similarity',
 'start',
 'subtree',
 'tensor',
 'text',
 'vector',
 'vocab']
```

You can identify the **root** of the sentence as determined by the dependency parser. Assuming the parsing is correct, this will be the main verb of the sentence.

```
print(sent.root)
```

narrantur

Each word in the sentence has an associated vector. Sentence (any **Span** in fact) has an associated vector as well that is the **mean of the vectors** of the words in the sentence. As this example uses the **lg** model, the vector has a length of 300.

```
print(sent.vector.shape)
```

(300,)

This vector then can be used to compute the similarity between two sentences. Here we see our example sentence compared to two related sentence: 1. a sentence where the character referred to is changed from Perseus to Ulysses; and 2. the active-verb version of the sentence.

```
sent.similarity(nlp('Haec narrantur a poetis de Ulixе.'))
```

0.9814933448498585

```
sent.similarity(nlp('Haec narrant poetae de Perseo.'))
```

0.7961655550941479

We can retrieve the start and end indices from the original document for each sentence.

```
sent_2 = sents[1]
start = sent_2.start
end = sent_2.end
print(start)
print(end)
print(sent_2.text)
print(doc[start:end].text)
```

7

15

Perseus filius erat Iovis, maximi deorum.

Perseus filius erat Iovis, maximi deorum.

References

SLP Chapter 2, Section 2.4.5 “Sentence Segmentation”, pp. 24 [link] (<https://web.stanford.edu/~jurafsky/slp3/>)
spaCy [SentenceRecognizer](#)

4 Word Tokenization

4.1 Word tokenization with LatinCy

```
# Imports & setup

import spacy
from pprint import pprint
nlp = spacy.load('la_core_web_md')
text = "Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius."
doc = nlp(text)
print(doc)
```

Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius.

Word tokenization is the task of splitting a text into words (and wordlike units like punctuation, numbers, etc.). For the LatinCy models, tokenization is the fundamental pipeline component on which all other components depend. SpaCy uses non-destructive, “canonical” tokenization, i.e. non-destructive, in that the original text can be untokenized, so to speak, based on `Token` annotations and canonical in that indices are assigned to each token during this process and these indices are used to refer to the tokens in other annotations. (Tokens *can* be separated or merged, but this requires the user to actively undo and redefine the tokenization output.) LatinCy uses a modified version of the default spaCy tokenizer that recognizes and splits enclitic *-que* using a rules-based process. (NB: It is in the LatinCy [development plan](#) to move enclitic splitting to a separate post-tokenization component.)

The spaCy Doc object is an iterable and tokens are the iteration unit.

```
tokens = [item for item in doc]
print(tokens)
```

[Haec, narrantur, a, poetis, de, Perseo, ., Perseus, filius, erat, Iovis, ,, maximi, deorum,

```
token = tokens[0]
print(type(token))
```

```
<class 'spacy.tokens.token.Token'>
```

The text content of a Token object can be retrieved with the `text` attribute.

```
for i, token in enumerate(tokens, 1):
    print(f'{i}: {token.text}')
```

```
1: Haec
2: narrantur
3: a
4: poetis
5: de
6: Perseo
7: .
8: Perseus
9: filius
10: erat
11: Iovis
12: ,
13: maximi
14: deorum
15: .
16: Avus
17: eius
18: Acrisius
19: appellabatur
20: .
```

Note again that the token itself is a spaCy Token object and that the `text` attribute returns a Python string even though their representations in the Jupyter Notebook look the same.

```
token = tokens[0]
print(f'{type(token)} -> {token}')
print(f'{type(token.text)} -> {token.text}')
```

```
<class 'spacy.tokens.token.Token'> -> Haec
<class 'str'> -> Haec
```

4.1.1 Token attributes and methods related to tokenization

Here are some attributes/methods available for spaCy `Token` objects that are relevant to word tokenization.

SpaCy keeps track of both the token indices and the character offsets within a doc using either the `i` or `idx` attributes, respectively...

```
print(token.doc)
```

Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius.

```
# token indices

for token in doc:
    print(f'{token.i}: {token.text}')
```

```
0: Haec
1: narrantur
2: a
3: poetis
4: de
5: Perseo
6: .
7: Perseus
8: filius
9: erat
10: Iovis
11: ,
12: maximi
13: deorum
14: .
15: Avus
16: eius
17: Acrisius
18: appellabatur
19: .
```

This is functionally equivalent to using `enumerate`...

```
# token indices, with enumerate

for i, token in enumerate(doc):
    print(f'{i}: {token.text}')
```

```
0: Haec
1: narrantur
2: a
3: poetis
4: de
5: Perseo
6: .
7: Perseus
8: filius
9: erat
10: Iovis
11: ,
12: maximi
13: deorum
14: .
15: Avus
16: eius
17: Acrisius
18: appellabatur
19: .
```

Another indexing option is the `idx` attribute which is the character offset of the token in the original Doc object.

```
# character offsets,
for token in doc:
    print(f'{token.idx}: {token.text}')
```

```
0: Haec
5: narrantur
15: a
17: poetis
24: de
27: Perseo
33: .
35: Perseus
```

```
43: filius
50: erat
55: Iovis
60: ,
62: maximi
69: deorum
75: .
77: Avus
82: eius
87: Acrisius
96: appellabatur
108: .
```

Observe these `idx` attributes relate to the character offsets from the original `Doc`. To illustrate the point, we will replace spaces with an underscore in the output. We can see from the output above that *narrantur* begins at `idx` 5 and that the next word *a* begins at `idx` 15. Yet *narrantur* is only 9 characters long and the difference between these two numbers is 10! This is because we need to account for whitespace in the original `Doc`. This is handled by the attribute `text_with_ws`.

```
print(doc.text[5:15].replace(' ', '_'))
```

```
narrantur_
```

```
print(f'text -> {doc[1].text} (length {len(doc[1].text)})')
print()
print(f'text_with_ws -> {doc[1].text_with_ws} (length {len(doc[1].text_with_ws)})')
```

```
text -> narrantur (length 9)
```

```
text_with_ws -> narrantur (length 10)
```

Accordingly, using the `text_with_ws` attribute (as opposed to simply the `text` attribute) we can reconstruct the original text. This is what was meant above by “non-destructive” tokenization. Look at the difference between a text joined using the `text` attribute and one joined using the `text_with_ws` attribute.

```
joined_tokens = ' '.join([token.text for token in doc])
print(joined_tokens)
print(joined_tokens == doc.text)

print()

reconstructed_text = ''.join([token.text_with_ws for token in doc])
print(reconstructed_text)
print(reconstructed_text == doc.text)
```

Haec narrantur a poetis de Perseo . Perseus filius erat Iovis , maximi deorum . Avus eius Acrisius
False

Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius
True

Because spaCy tokenization is set from the outset, you can traverse the tokens in a `Doc` objects from the tokens themselves using the `nbord` method. This method takes an integer argument that specifies the number of tokens to traverse. A positive integer traverses the tokens to the right, a negative integer traverses the tokens to the left.

```
print(doc[:6])
print('-----')
print(f'{doc[3]}, i.e. i = 3')
print(f'{doc[3].nbord(-1)}, i.e. i - 1 = 2')
print(f'{doc[3].nbord(-2)}, i.e. i - 2 = 1')
print(f'{doc[3].nbord(1)}, i.e. i + 1 = 4')
print(f'{doc[3].nbord(2)}, i.e. i + 2 = 5')
```

Haec narrantur a poetis de Perseo

poetis, i.e. i = 3
a, i.e. i - 1 = 2
narrantur, i.e. i - 2 = 1
de, i.e. i + 1 = 4
Perseo, i.e. i + 2 = 5

4.1.2 Customization of the spaCy tokenizer in LatinCy

LatinCy aims to tokenize the *que* enclitic in Latin texts. As noted above this is currently done through a rule-based approach. Here is the custom tokenizer code (beginning at this line in

the [code](#)) followed by a description of the process. Note that this process is based on the following recommendations in the spaCy documentation: <https://spacy.io/usage/training#custom-tokenizer>.

```
from spacy.util import registry, compile_suffix_regex

@registry.callbacks("customize_tokenizer")
def make_customize_tokenizer():
    def customize_tokenizer(nlp):
        suffixes = nlp.Defaults.suffixes + [
            "que",
            "qve",
        ]
        suffix_regex = compile_suffix_regex(suffixes)
        nlp.tokenizer.suffix_search = suffix_regex.search

    for item in que_exceptions:
        nlp.tokenizer.add_special_case(item, [{"ORTH": item}])
        nlp.tokenizer.add_special_case(item.lower(), [{"ORTH": item.lower()}])
        nlp.tokenizer.add_special_case(item.title(), [{"ORTH": item.title()}])
        nlp.tokenizer.add_special_case(item.upper(), [{"ORTH": item.upper()}])

    return customize_tokenizer
```

Basically, we treat *que* (and its case and u/v norm variants) as punctuation. These are added to the `Defaults.suffixes`. If no other intervention were made, then any word ending in *que* or a variant would be split into a before-*que* part and *que*. Since there are large number of relatively predictable words that end in *que*, these are maintained in a list called `que_exceptions`. All of the words in the `que_exceptions` list are added as a “special case” using the tokenizer’s `add_special_case` method and so will not be split. The `que_exceptions` lists is as follows:

```
que_exceptions = ['quisque', 'quidque', 'quicque', 'quodque', 'cuiusque', 'cuique', 'quemque']
```

You can see these words in the `rules` attribute of the tokenizer.

```
# Sample of 10 que rules from the custom tokenizer

tokenizer_rules = nlp.tokenizer.rules
print(sorted(list(set([rule.lower() for rule in tokenizer_rules if 'que' in rule]))[:10]))
```

['absque', 'abusque', 'adaeque', 'adusque', 'aeque', 'antique', 'atque', 'circumundique', 'c']

With the exception of basic enclitic splitting, the LatinCy tokenizer is the same as the default spaCy tokenizer. The default spaCy tokenizer is described in detail in the [spaCy documentation](#). Here are some useful attributes/methods for working with LatinCy.

Tokenize a string without any other pipeline annotations with a call.

```
tokens = nlp.tokenizer(text)
print(tokens)
print(tokens[0].text)
print(tokens[0].lemma_) # Note that there is no annotation here because since the tokenizer
```

Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acris.
Haec

A list of texts can be tokenized in one pass with the `pipe` method. This yields a generator object where each item is Doc object of tokenized-only texts

```
texts = ["Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum; avus e"]
tokens = list(nlp.tokenizer.pipe(texts))

print(len(tokens)) # number of documents
print(len(tokens[0])) # number of tokens in first document
```

3
76

You can get an explanation of the tokenization “decisions” using the `explain` method. In the example below, we see how the *que* in *virumque* is treated as a suffix (as discussed above) and so is split during tokenization.

```
tok_exp = nlp.tokenizer.explain('arma virumque cano')
print(tok_exp)
```

```
[('TOKEN', 'arma'), ('TOKEN', 'virum'), ('SUFFIX', 'que'), ('TOKEN', 'cano')]
```

```
tokens = nlp.tokenizer('arma uirumque cano')
for i, token in enumerate(tokens):
    print(f'{i}: {token.text}')
```


0: arma
1: uirum
2: que
3: cano

References

SLP Chapter 2, Section 2.4.2 “Word Tokenization”, pp. 18-20 [link](#)

5 Lemmatization

5.1 Lemmatization with LatinCy

Lemmatization is the task of mapping a token in a text to its dictionary headword. With the default LatinCy pipelines, two components are used to perform this task: 1. spaCy’s [Edit Tree Lemmatizer](#) and 2. a second custom `Lookup Lemmatizer`, named in the pipeline “trainable_lemmatizer” and “lookup_lemmatizer” respectively. In the first lemmatization pass, a probabilistic tree model is used to predict the transformation from the token form to its lemma. A second pass is made at the end of the pipeline which checks the token form against a large (~1M item) lemma dictionary (i.e. lookups) for ostensibly unambiguous forms; if a match is found, the lemma is overwritten with the corresponding value from lookup. The two-pass logic largely follows the approach recommended in Burns (2018) and Burns (2020), as facilitated by the spaCy pipeline architecture.

```
# Imports & setup

import spacy
from pprint import pprint
nlp = spacy.load('la_core_web_sm')
text = "Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius."
doc = nlp(text)
print(doc)
```

Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acrisius.

Note here the two lemmatizer components that are included in the pipeline, i.e. “trainable_lemmatizer” and “lookup_lemmatizer”...

```
print(nlp.pipe_names)
```

```
['senter', 'normer', 'tok2vec', 'tagger', 'morphologizer', 'trainable_lemmatizer', 'parser',
```

Once a text is annotated using the LatinCy pipeline, i.e. as part of the Doc creation process, lemmas can be found as annotations of the `Token` objects...

```
sample_token = doc[0]

print(f'Sample token: {sample_token.text}')
print(f'Sample lemma: {sample_token.lemma_}')
```

```
Sample token: Haec
Sample lemma: hic
```

```
import tabulate

data = []

tokens = [item for item in doc]

for token in tokens:
    data.append([token.text, token.lemma_])

print(tabulate.tabulate(data, headers=['Text', 'Lemma']))
```

Text	Lemma
Haec	hic
narrantur	narro
a	ab
poetis	poeta
de	de
Perseo	Perseo
.	.
Perseus	Perseus
filius	filius
erat	sum
Iovis	Iuppiter
,	,
maximi	
deorum	deus
.	.
Avus	auus
eius	is

Acrisius	Acrisius
appellabatur	appello
.	.

The `lemma_` attribute has the type `str` and so is compatible with all string operations...

```
print(f'Token: {tokens[0].text}')
print(f'Lemma: {tokens[0].lemma_}')
print(f'Lowercase lemma: {tokens[0].lemma_.lower()}')
```

```
Token: Haec
Lemma: hic
Lowercase lemma: hic
```

The `lemma_` attribute, though, is only the human-readable version of the lemma. Internally, spaCy uses a hash value to represent the lemma, which is stored in the `lemma` attribute. (Note the lack of trailing underscore.)

```
print(f'Token: {tokens[1].text}')
print(f'Human-readable lemma: {tokens[1].lemma_}')
print(f'spaCy lemma key: {tokens[1].lemma}')
```

```
Token: narrantur
Human-readable lemma: narro
spaCy lemma key: 11361982710182407617
```

In order to compare the two different lemmatization passes, we can create two copies of the LatinCy pipeline, each with one of the two lemmatizers removed...

```
import copy

P1 = copy.deepcopy(nlp)
P1.disable_pipes(['tagger', 'morphologizer', 'lookup_lemmatizer'])
print(f'First pipeline components: {P1.pipe_names}')

P2 = copy.deepcopy(nlp)
P2.disable_pipes(['tagger', 'morphologizer', 'trainable_lemmatizer'])
print(f'Second pipeline components: {P2.pipe_names}')
```

```
First pipeline components: ['sender', 'normer', 'tok2vec', 'trainable_lemmatizer', 'parser',
Second pipeline components: ['sender', 'normer', 'tok2vec', 'parser', 'lookup_lemmatizer', '']
```

We can then run the same text through both pipelines and compare the results side-by-side...

```
P1_annotations = P1(text)
P2_annotations = P2(text)

data = []

for p1_token, p2_token in zip(P1_annotations, P2_annotations):
    data.append([p1_token.text, p1_token.lemma_, p2_token.lemma_])

print(tabulate.tabulate(data, headers=['Text', 'Trainable lemmatizer', 'Lookup lemmatizer']))
```

Text	Trainable lemmatizer	Lookup lemmatizer
-----	-----	-----
Haec	hic	
narrantur	narro	narro
a	ab	
poetis	poetus	poeta
de	de	
Perseo	Perseo	
.	.	.
Perseus	Perseus	
filius	filius	filius
erat	sum	sum
Iovis	Iovis	Iuppiter
,	,	,
maximi		
deorum	deus	deus
.	.	.
Avus	avus	auus
eius	is	is
Acrisius	Acrisius	
appellabatur	appello	appello
.	.	.

Note specifically the lemmatization of *Iovis*—since it is a highly irregular form, it is not surprising that the Edit Tree Lemmatizer has manufactured a potential, but incorrect, lemma based on the root *Iov-*. Since *Iovis* is an unambiguous Latin form and has been added to the LatinCy lookups, the Lookup Lemmatizer steps in to correct the erroneous first pass. The lookup data can be found a custom form of the spaCy `spacy-lookups-data` package [https://github.com/diyclassics/spacy-lookups-data/tree/master/spacy_lookups_data/data]. These lookups are installed as a dependency of each of the LatinCy pipelines.

The code below shows how to access the lookup data directly...

```
# Load the lookups data

from spacy.lookups import Lookups, load_lookups

blank_nlp = spacy.blank("la")
lookups = Lookups()

lookups_data = load_lookups(lang=blank_nlp.vocab.lang, tables=["lemma_lookup"])
LOOKUPS = lookups_data.get_table("lemma_lookup")

print(LOOKUPS['Iovis'])
```

Iuppiter

References

NLTK Chapter 3, Section 3.6 “Normalizing text” [link](#)

SLP Chapter 2, Section 2.6 “Word normalization, lemmatization and stemming”, pp. 23-24 [link](#)

spaCy EditTreeLemmatizer [link](#)

6 POS Tagging

6.1 POS Tagging with LatinCy

Part-of-speech tagging is the task of mapping a token in a text to its part of speech, whether ‘noun’ or ‘verb’ or ‘prepositon’ and so on. There are two components in the default LatinCy pipelines that provide such annotations, that is the `tagger` and the `morphologizer`. Ostensibly, the `tagger` provides language-specific, fine-grain POS tags and the `morphologizer` provides coarse-grain tags (as defined by the UD [Universal POS tags](#)); at present, the LatinCy models have a high degree of overlap between these two tagsets and there are effectively *no* fine-grain tags.

```
# Imports & setup

import spacy
from pprint import pprint
nlp = spacy.load('la_core_web_lg')
text = "Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus ei
doc = nlp(text)
print(doc)
```

Haec narrantur a poetis de Perseo. Perseus filius erat Iovis, maximi deorum. Avus eius Acris

Note here the two “tagging” components that are included in the pipeline, i.e. “tagger” and “morphologizer”...

```
print(nlp.pipe_names)
```

```
['senter', 'normer', 'tok2vec', 'tagger', 'morphologizer', 'trainable_lemmatizer', 'parser',
```

Once a text is annotated using the LatinCy pipeline, i.e. as part of the `Doc` creation process, tags can be found as annotations of the `Token` objects. The coarse-grain tags are stored in the `pos_` attribute; fine-grain tags are stored in the `tag_` attribute.

```

sample_token = doc[1]

print(f'Sample token: {sample_token.text}')
print(f'Sample POS: {sample_token.pos_}')
print(f'Sample TAG: {sample_token.tag_}')

```

```

Sample token: narrantur
Sample POS: VERB
Sample TAG: verb

```

Note the high degree of overlap between the coarse-grain and fine-grain tags in the LatinCy models in the chart below. That said, it is perhaps worth paying more attention to where the tagsets do not overlap. In the LatinCy training, conventions for classes of words to be labeled, say, “DET” (as in the *Haec pos_*) or “AUX” (as in *est*) are inferred from usage in the six different treebanks used. A sense of the inconsistency in the tagsets can be gleaned from the following page: (<https://universaldependencies.org/la/>); note also the important work of Gamba and Zeman (2023) in this area.

```

import tabulate

data = []

tokens = [item for item in doc]

for token in tokens:
    data.append([token.text, token.pos_, token.tag_])

print(tabulate.tabulate(data, headers=['Text', "POS", "TAG"]))

```

Text	POS	TAG
-----	-----	-----
Haec	DET	pronoun
narrantur	VERB	verb
a	ADP	preposition
poetis	NOUN	noun
de	ADP	preposition
Perseo	PROPN	proper_noun
.	PUNCT	punc
Perseus	PROPN	proper_noun
filius	NOUN	noun
erat	AUX	verb

Iovis	PROPN	proper_noun
,	PUNCT	punc
maximi	ADJ	adjective
deorum	NOUN	noun
.	PUNCT	punc
Avus	NOUN	noun
eius	PRON	pronoun
Acrisius	PROPN	proper_noun
appellabatur	VERB	verb
.	PUNCT	punc

As with the lemma annotations, the `pos_` and `tag_` attributes are only the human-readable of the lemma. Internally, spaCy uses a hash value to represent this, again noting the lack of trailing underscore...

```
print(f"Token: {tokens[1].text}")
print(f'Human-readable TAG: {tokens[1].tag_}')
print(f'spaCy TAG key: {tokens[1].tag}')
```

```
Token: narrantur
Human-readable TAG: verb
spaCy TAG key: 6360137228241296794
```

This tag 'key' can be looked up in spaCy's `NLP.vocab.strings` attribute...

```
T = nlp.get_pipe('tagger')
tag_lookup = T.vocab.strings[6360137228241296794] # also would work on `nlp`, i.e. nlp.vocab

print(f'TAG key: {tokens[1].tag}')
print(f'Human-readable TAG: {tag_lookup}')
```

```
TAG key: 6360137228241296794
Human-readable TAG: verb
```

The same process applies to the POS 'keys'...

```
print(f'Token: {tokens[1].text}')
print(f'Human-readable POS: {tokens[1].pos_}')
print(f'spaCy POS key: {tokens[1].pos}')
```

```
Token: narrantur
Human-readable POS: VERB
spaCy POS key: 100
```

```
M = nlp.get_pipe("morphologizer")
pos_lookup = M.vocab.strings[100]
print(f'POS key: {tokens[1].pos}')
print(f'Human-readable POS: {pos_lookup}')
```

```
POS key: 100
Human-readable POS: VERB
```

We can use the `label_data` attribute from the `morphologizer` component to derive the complete (coarse-grain) tagset...

```
def split_pos(morph):
    if 'POS=' in morph:
        return morph.split('POS=')[1].split('|')[0]
    else:
        return None

tagset = sorted(list(set([split_pos(k) for k, v in M.label_data['morph'].items() if split_pos(k)])))
print(tagset)
```

```
['ADJ', 'ADP', 'ADV', 'AUX', 'CCONJ', 'DET', 'INTJ', 'NOUN', 'NUM', 'PART', 'PRON', 'PROPN',
```

SpaCy has an `explain` methods that can show human-readable descriptions of these standard tags...

```
data = []

for tag in tagset:
    data.append([tag, spacy.explain(tag)])

print(tabulate.tabulate(data, headers=['TAG', 'Description']))
```

TAG	Description
ADJ	adjective
ADP	adposition

ADV	adverb
AUX	auxiliary
CCONJ	coordinating conjunction
DET	determiner
INTJ	interjection
NOUN	noun
NUM	numeral
PART	particle
PRON	pronoun
PROPN	proper noun
PUNCT	punctuation
SCONJ	subordinating conjunction
VERB	verb
X	other

It may also be useful to know the “confidence” of the tagger in making its decision. We can derive this from the output tagger’s `model.predict` method. This returns (at least, in part) a ranked list of per-scores, the maximum value of which determines the final annotation.

```
# Helper function to get tagging scores

def get_tagging_scores(doc, n=3):
    # cf. https://stackoverflow.com/a/69228515
    scores = []
    tagger = nlp.get_pipe('tagger')
    labels = tagger.labels
    for token in doc:
        token_scores = tagger.model.predict([doc])[0][token.i]
        r = [*enumerate(token_scores)]
        r.sort(key=lambda x: x[1], reverse=True)
        scores.append([(labels[i], p) for i, p in r[:n]])
    return scores
```

```
# Get the top 3 tags by score for each token in the Doc

tagging_probs = get_tagging_scores(doc)

for token in doc:
    print(f'Token: {token.text}', end='\n\n')
    data = []
    for label, prob in tagging_probs[token.i]:
        data.append([label, prob])
```

```
print(tabulate.tabulate(data, headers=['Label', 'Score']))
break
```

Token: Haec

Label	Score
pronoun	10.7023
adjective	8.26614
noun	3.13534

References

NLTK Chapter 5 “Categorizing and tagging words” [link](#)

SLP Chapter 8 “Sequence labeling for parts of speech and named entities” [link](#)

spaCy [Tagger](#) and [Morphologizer](#)

Part IV

spaCy Tasks

Open LatinCy Projects

Listed below are some open projects that I would like to see implemented in the LatinCy pipelines. If you are interested in contributing to the development of these projects or have recommendations for additional projects that could be added, please contact me through the LatinCy GitHub or HuggingFace repositories.

Model-based enclitic splitting

As discussed in the [Word Tokenization](#) chapter, enclitic splitting is currently limited to *que* (and variants) and is a rules-based process that is part of the LatinCy custom tokenizer. It would be preferable to make enclitic splitting a separate pipeline component and moreover one that is model-based. This component could be placed immediately after the tokenizer and use the [Retokenizer.split method](#) to reconstruct token sequences where valid enclitics are identified. This would have the added advantage of allowing enclitic splitting to be “turned off”, so to speak, by removing the component from the pipeline.

References

- Bird, Steven, Ewan Klein, and Edward Loper. 2015. *Natural Language Processing with Python: Analyzing Text with the Natural Language Toolkit*. 2nd Edition. <https://www.nltk.org/book/>.
- Burns, Patrick J. 2018. “Backoff Lemmatization as a Philological Method.” In *Digital Humanities 2018, DH 2018, Book of Abstracts, El Colegio de México, UNAM, and RedHD, Mexico City, Mexico, June 26-29, 2018*.
- . 2020. “Ensemble Lemmatization with the Classical Language Toolkit.” *Studi e Saggi Linguistici* 58 (1): 157–76. <https://doi.org/10.4454/ssl.v58i1.273>.
- Clayman, Dee. 1981. “Sentence Length in Greek Hexameter Poetry.” *Quantitative Linguistics* 11: 107–36. <https://papers.ssrn.com/abstract=1627358>.
- Gamba, Federica, and Daniel Zeman. 2023. “Universalising Latin Universal Dependencies: A Harmonisation of Latin Treebanks in UD.” In *Proceedings of the Sixth Workshop on Universal Dependencies (UDW, GURT/SyntaxFest 2023)*, edited by Loïc Grobol and Francis Tyers, 7–16. Washington, D.C.: Association for Computational Linguistics. <https://aclanthology.org/2023.udw-1.2>.
- Janson, Tore. 1964. “The Problems of Measuring Sentence-Length in Classical Texts.” *Studia Linguistica* 18 (1): 26–36. <https://doi.org/10.1111/j.1467-9582.1964.tb00443.x>.
- Jurafsky, Daniel, and James H. Martin. 2020. “Speech and Language Processing (3rd Edition, Draft).” <https://web.stanford.edu/~jurafsky/slp3/>.
- Wake, William C. 1957. “Sentence-Length Distributions of Greek Authors.” *Journal of the Royal Statistical Society. Series A (General)* 120 (3): 331–46. <https://www.jstor.org/stable/2343104>.